

Architecture Review — Critical Questions

Event Finder — Architecture Review: Critical Questions

Strategy & Scope

- 1) What is the real target scale: dozens, hundreds, or thousands of events, and how frequently does the catalog refresh?
- 2) Who owns the ingestion sources and what legal/usage constraints apply to each source?
- 3) Are events local-only or multi-region? If multi-region, do we need locale/region awareness (time zones, languages, date formats)?
- 4) What is the expected user journey: quick lookup or browsing/discovery? Any mobile-first constraints?

Data Quality & Governance

- 5) How do we handle duplicates across sources, conflicting fields (e.g., venue names), and missing data (time, location)?
- 6) Do we need canonicalization rules (e.g., normalize dates, locations, and organizer names)?
- 7) Should we track data lineage to know which source produced each field?

Search & Relevance

- 8) What are the matching expectations (exact match, stemming, synonyms, typo tolerance)?
- 9) Do we need ranking signals beyond text match (recency, popularity, distance, organizer reputation)?
- 10) Should we support structured filters (date range, location radius, price, category) in the near term?

Experience & Content

- 11) Do we need result cards with fielded data or is the narrative response the primary format? Both?
- 12) Should the UI suggest related queries or clarifications when zero results are found?
- 13) Accessibility targets (keyboard navigation, screen readers, contrast)?

Operations & Reliability

- 14) What uptime/latency SLOs matter (if any)? Is this strictly best-effort or user-facing with expectations?

- 15) Do we require persistence across restarts now, or is ephemeral acceptable until phase 2?
- 16) What's the update mechanism for ingestion: manual trigger, scheduled job, or webhook-driven?
- 17) Do we need observability basics now (request logs, ingestion metrics, error tracking)?

Security & Compliance

- 18) Will this be exposed publicly? If yes, when do we add TLS, rate limiting, and abuse controls?
- 19) Any PII or sensitive fields planned now or later (e.g., RSVP emails)?
- 20) What is the threat model for data poisoning from untrusted sources?

Roadmap & Governance

- 21) Who prioritizes the roadmap (search quality vs. ingestion coverage vs. UI polish)?
- 22) What are success metrics (conversion to event page, time to first useful result, query success rate)?
- 23) Do we need an admin view for curation, audit logs, and content takedowns?

Event Finder — High-Level Architecture (Revised)

Event Finder — High-Level Architecture (Revised)

1) Overview & Objectives

A pragmatic event discovery service that ingests event listings from approved sources, normalizes them, provides fast search, and returns clear summaries in a simple web UI. Priorities: correctness, clarity, and low operational overhead.

2) Capabilities

- Ingestion: Collect and normalize events from configured sources (HTML, feeds, files).
- Catalog: Maintain a clean, deduplicated, structured store of events.
- Search & Relevance: Text search with sensible ordering (recency-first), optional structured filters.
- Response Generation: Human-friendly narratives and optional card-style results.
- Presentation: Minimal, responsive web interface with a search box and results.

3) Context (Conceptual Flow)

Sources → Ingestion → Normalization/Deduplication → Catalog → Search & Relevance → Response → Web UI

4) Logical Components

- Source Connectors: One per source type (HTML/feeds/files); encapsulate parsing and error handling.
- Normalizer/Deduper: Standardizes fields (name, date/time, location), resolves duplicates via heuristics.
- Catalog Service: Owns the event model; supports upsert, list, and read for search.
- Search Service: Provides keyword search and optional filters (date range, location, category).
- Response Service: Formats results as narratives and/or structured cards.
- Web UI: Search input, loading state, results; zero-login, mobile-friendly.
- Admin (Future): Manual review, source health, re-ingestion triggers.

5) Data Model (Business-Level)

Event

- Core: name, date, time, location (venue, city/region), description

- Optional (future): category/tags, organizer, price, URL, image, geo-coordinates, source id
- Metadata: source, ingestion timestamp, normalization flags

6) Key Cross-Cutting Concerns

- Data Quality: Normalization rules, duplicate detection, missing-field fallbacks.
- Relevance: Recency bias by default; optional distance/categorization in phase 2.
- Observability: Basic logs for ingestion and search; error summaries.
- Security & Legal: Use only approved sources; do not expose admin/debug endpoints publicly.
- Privacy: Avoid PII unless explicitly required and governed.
- Accessibility: Basic a11y (labels, semantic markup, keyboard navigation).

7) Interaction Scenarios

Search (Happy Path)

1. User enters a query in the UI.
2. UI calls Search Service with query and optional filters.
3. Search returns candidate events from the Catalog.
4. Response Service formats results as a short narrative (+ optional cards).
5. UI renders results and suggests refinements if needed.

Ingestion/Refresh

1. Admin or schedule initiates ingestion for configured sources.
2. Connectors fetch data; Normalizer/Deduper standardizes and merges.
3. Catalog is updated with upserts; Search index reflects updates.
4. Logs and summaries indicate freshness and failures.

8) Non-Functional Goals (Target)

- Simplicity: Minimal moving parts; single small host acceptable.
- Performance: Snappy results for small-to-medium catalogs (sub-second typical).
- Resilience: Graceful behavior with missing or malformed source data.
- Portability: Able to run locally and in a small cloud instance.
- Extensibility: Pluggable connectors, fields, and filters.

9) Operational View

- Environments: Local dev, optional small cloud.
- Configuration: Source list, refresh cadence, feature flags (filters on/off).
- Deployment: App server with static UI; no public debug endpoints.
- Monitoring: Request logs, ingestion counts, error rates; simple dashboards later.
- Runbooks: “Search is slow” and “Source is failing” quick checks.

10) Risks & Mitigations

- Source Volatility: Formats drift — isolate parsers and add validation.
- Data Drift/Noise: Duplicates, inconsistent fields — normalization & merge heuristics.
- Exposure Risk: Public debug or open endpoints — disable debug, add basic protections before public launch.
- Relevance Gaps: Exact-match brittleness — plan synonyms/typo tolerance roadmap.
- Scalability Ceiling: Growth beyond a single host — plan step-up (persistent store, external search).

11) Roadmap

Phase 0 (Now): MVP with keyword search, narrative responses, seed data, minimal UI.

Phase 1: Persistent catalog, basic filters (date range, city/region), admin trigger for refresh.

Phase 2: Relevance tuning (recency, synonyms), multi-source connectors, basic metrics dashboard.

Phase 3: Geo-awareness, richer cards (images, URLs), rate limiting, TLS, and basic auth for admin.

12) Success Metrics

- Query Success Rate (non-empty useful results)
- Time to First Result (p50/p90)
- Data Freshness (time since last successful ingestion)
- Zero-Result Recovery (suggestions clicked)
- Operational Health (ingestion error rate)